

Run Time & Performance

Kurt Schmidt

Dept. of Computer Science, Drexel University

June 5, 2017

Examples are taken from Kernighan & Pike, *The Practice of Programming*, Addison-Wesley, 1999



Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Intro

Performance

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Objective: To learn when and how to optimize the performance of a program.

The first principle of optimisation is don't.

- Knowing how a program will be used, and the environment it runs in, is there any benefit to making it faster?
- Which areas of the program should we focus on?

Strategy

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- Use the simplest, cleanest algorithms and data structures appropriate for the task
- Enable compiler options to generate the fastest possible code
 - Modern compilers optimise by default
 - Modern compilers are very good at their jobs
- Then, measure performance to see if changes are needed

Strategy (cont.)

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- Assess what changes to the program will have the most effect
 - Use a profiler to find hotspots in your code
- Make changes incrementally, re-assess
 - Consider alternative algorithms
 - Tune the code
 - Consider a lower-level language
 - Maybe just for time-sensitive components
 - Always retest your code!

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Profiling Code

Profiler – gprof

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- A profiler watches your program run
- Reports back a bunch of information, including
 - How much time was spent in each function
 - How many times each function was called
- Use this information to find bottlenecks (areas worth improvement) in your code
- Profilers exist for most common languages, including Java, Python, and Haskell
- We will use `gprof`, which can be run with programs compiled with a `gnu` compiler

Using gprof

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- Use the `-p` option to compile extra information into your program:

```
$ gcc -p driver.c quicksort.c -o mySort
```

- Now run the program once, to generate metrics:

```
$ ./mySort < ins.10000 > /dev/null
```

- Note, a new data file has appeared:¹

```
$ ls -ot | head -n3
total 3864
-rw-r--r-- 1 kschmidt 747 Aug 4 16:05 gmon.out
-rwxrwxr-x 1 kschmidt 9102 Aug 4 16:05 mySort
```

¹It's raw data. Don't look at it yet

Reading the gprof Report

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- Supply gprof w/the name of the executable to see report on stdout:

```
$ gprof mySort
```

- Report contains 2 tables of data, with description of the information:

```
sample
=====
sample Each sample counts as 0.01 seconds.
sample
sample % cumulative self          self total
sample time  seconds seconds    calls us/call us/call name
sample 61.41    0.25    0.25      500  503.60  805.76 quicksort
sample 36.85    0.40    0.15 45721062    0.00    0.00 swap
sample  1.23    0.41    0.01
sample
sample ...
sample
=====
```

Reading the gprof Report (cont.)

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

```
...
granularity: each sample hit covers 2 byte(s) for 2.45% of 0.41 sec

index % time    self children   called    name
-----
[1]    100.0    0.01    0.40           <spontaneous>
           0.25    0.15    500/500        main [1]
           quicksort [2]
-----
           66666956        quicksort [2]
           0.25    0.15    500/500        main [1]
[2]    98.8    0.25    0.15    500+66666956 quicksort [2]
           0.15    0.00 45721062/45721062 swap [3]
           66666956        quicksort [2]
-----
           0.15    0.00 45721062/45721062 quicksort [2]
[3]    37.0    0.15    0.00 45721062        swap [3]
-----
...
```

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

**Run-time
Analysis**

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Run-time Analysis

Asymptotic Run Time

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- We would like to be able to compare algorithms, for arbitrarily large inputs
- We want to evaluate the growth of algorithms vs. input
 - We don't care about, e.g., processor speed, the leading coefficient, nor lower-order terms
 - E.g., linear search. If the array size doubles, so does the run-time $\implies \Theta(n)$
 - Selection sort is a quadratic sort, $\Theta(n^2) \implies$ doubling input size will quadruple run time (for large n)
 - Accessing an element of an array is a constant-time operation, $\Theta(1)$, regardless of size of array

Lower Bound (Loose) – Little Omega

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Consider function $T_n = 5n^2 + 17n - 12$

$$\lim_{n \rightarrow \infty} \frac{5n^2 + 17n - 12}{n} = \infty \implies T_n \in \omega(n)$$

- We say T_n is *bound below* (loosely) by n
- We say T_n grows *strictly faster* (asymptotically) than n
- T_n can not be bound above by n

More generally:

$$\lim_{n \rightarrow \infty} \frac{T_n}{f_n} = \infty \implies T_n \in \omega(f_n)$$

Upper Bound (Loose)– Little Oh

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

$$\lim_{n \rightarrow \infty} \frac{5n^2 + 17n - 12}{n^3} = 0 \implies T_n \in o(n^3)$$

- We say T_n is *bound above* (loosely) by n^3
- We say T_n grows *strictly more slowly* (asymptotically) than n^3
- T_n can not be bound below by n^3

More generally:

$$\lim_{n \rightarrow \infty} \frac{T_n}{f_n} = 0 \implies T_n \in o(f_n)$$

Asymptotically Equivalent – Theta

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

$$\lim_{n \rightarrow \infty} \frac{5n^2 + 17n - 12}{n^2} = 5 \implies T_n \in \Theta(n^2)$$

- We say T_n grows like n^2 (asymptotically)
- T_n can bound below, and above, by n^2

More generally:

$$\lim_{n \rightarrow \infty} \frac{T_n}{f_n} = c, c \in \mathbb{R}^+ \implies T_n \in \Theta(f_n)$$

Note: This does not mean that T_n is polynomial

Notes on Asymptotic Equivalence

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- Two equivalent functions needn't look the same
- Just grow similarly
- Consider $y = x + \sin x$
 - y grows like a line
 - Bound above by $y = x + 1$
 - Bound below by $y = x - 1$
 - Clearly not a line

Quick Observations

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

$$f \in \mathbf{o}(g) \iff g \in \omega(f)$$

$$T \in \Theta(f) \implies T \notin \mathbf{o}(f)$$

$$T \in \Theta(f) \implies T \notin \omega(f)$$

$$T \in \omega(f) \implies T \notin \Theta(f)$$

$$T \in \mathbf{o}(f) \implies T \notin \Theta(f)$$

Big-Oh, -Omega

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Upper (lower) bound, may or may not be tight

$$O(f) = o(f) \cup \Theta(f)$$

$$\Omega(f) = \omega(f) \cup \Theta(f)$$

We have these observations:

$$o(f) \subset O(f)$$

$$\omega(f) \subset \Omega(f)$$

Finally,

$$T \in \Theta(f) \iff T \in O(f) \wedge T \in \Omega(f)$$

Qualitative Statements

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

T is $O(f)$

- “ T grows no faster than f ”
- “ T is bound above by f ”
- “ f is an upper bound for T ”

T is $\Omega(f)$

- “ T grows no slower than f ”
- “ T is bound below by f ”
- “ f is a lower bound for T ”

Ranking of Common Functions

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

For reference, here are some common functions, in increasing order:

1 (constant)

n^2

$\log n$

\vdots

\sqrt{n}

n^p

n

c^n

$n \log n$

$n^n \approx n!$

$n\sqrt{n}$

Note, $\log n \in o(n^p), p \in \mathbb{R}, \forall p > 0$

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Timing Programs

Timing – time

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- If we can't evaluate the algorithm, we can run the program with various inputs, time each run
- Various languages may provide their own mechanism for timing from within the program
- The `time` utility takes a program, with arguments, to run
 - You now know why you type `date` to see what the time is

Using the `time` Utility

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

```
time [options] cmd [cmd_args]
```

options Options to modify behavior of `time`. Must precede `cmd`

cmd The program run you want to time

cmd_args Arguments, including options, to be passed to `cmd`

Note: This is a built-in in Bash, Tenex C-Shell, and others.

time example

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

From our previous example:

```
$ time ./mySort < ins.10000 > /dev/null
```

- Output to screen is expensive
- We're not interested in that time

Output from time (to stderr):

```
real 0m7.572s
user 0m7.555s
sys 0m0.004s
```


Description of Output

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

real The wall clock. Total time elapsed. Keeps ticking, even if your program is sliced out

user The actual time your program spent running, in user mode

sys The actual time your program spent running, in kernel mode

- The sum of the user and sys times is probably what you want

Using Bash's `time` in a Script

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- The `-p` option be handier for parsing:

```
$ /usr/bin/time ./mySort < ins.10000 &> /dev/null
```

```
real 7.572
user 7.555
sys 0.004
```

- Storing it in a variable takes some contortions:

```
$ result=$( { time -p ./$prog <ins.10000 >/dev/null ; } 2>&1 )
$ echo $result
```

```
real 7.572 user 7.555 sys 0.004
```

time Built-in v. Utility

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- There is a utility (not a shell built-in)
 - On tux, installed as `/usr/bin/time`
 - Can also report on other metrics
 - Output values and format can be customised (see `-f`)
- Bash, Tenex C Shell, and others have a built-in `time` command, which doesn't allow for customising the output format
 - The built-ins are a bit slicker, parsing up the command line
 - For e.g., the following produces no output (why?), but works fine w/the shell built-in

```
$ /usr/bin/time ./mySort < ins.10000 &> /dev/null
```

Alternatives to Timing Program

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- Computers have gotten fast
 - Makes it a little harder to grab good numbers
- Alternatively, we could be creative, use metrics from a profiler
 - E.g., we could count the number of calls to `swap` for various sized inputs to our sort
 - Or, we might use a function to compare elements in a sort, use a profiler to count the number of times items are compared
- We can use this data in the same way, plot it, see how it grows

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

**Program
Growth**

Notes on
Scale

Program Growth

Estimating Program Growth

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- Run your program on various-sized inputs, collect times
 - Get a good number of points
 - Discard very small results
 - Provide inputs large enough to get past lower-order noise
- Find an upper and lower bound
- Consider $\frac{T_n}{f_n}$ for various functions f
 - Identify upper and lower bounds
 - Try to pinch in, get the upper and lower bounds closer to each other
 - You might not get them to meet

Estimating Program Growth – e.g.

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Consider times T discovered for various input sizes n :

n	$T(n)$
10	3908.51
20	20657.40
30	55954.53
40	113992.17
50	198284.36
60	311920.28
70	457689.75
80	638156.74
90	855706.82
100	1112580.00

■ $T(n)$ appears to be increasing w/out bound

■ So, T is not constant

■ Maybe. We don't *know* this

Let's compare to $f(n) = n$

E.g. – Line is a Lower Bound

Run Time & Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

n	$T(n)/n$
10	390.85
20	1032.87
30	1865.15
40	2849.80
50	3965.69
60	5198.67
70	6538.43
80	7976.96
90	9507.85
100	11125.80

- $T(n)/n$ also appears to be increasing, w/out bound

- So, $f(n) = n$ looks like a lower bound

- I.e., $T(n) \in \Omega(n)$

- If not tight, if it increases w/out bound, then $T(n) \in \omega(n)$

Let's try $f(n) = n^2$:

E.g. – Quadratic is a Lower Bound

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

n	$T(n)/n^2$
10	39.09
20	51.64
30	62.17
40	71.25
50	79.31
60	86.64
70	93.41
80	99.71
90	105.64
100	111.26

- $T(n)/n^2$ also appears to be increasing

- $f(n) = n^2$ looks like a lower bound

- i.e., $T(n) \in \Omega(n^2)$

- If not tight, if it increases w/out bound, then $T(n) \in \omega(n^2)$

Let's try $f(n) = n^3$:

E.g. – Cubic is an Upper Bound

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

n	$T(n)/n^3$
10	3.91
20	2.58
30	2.07
40	1.78
50	1.59
60	1.44
70	1.33
80	1.25
90	1.17
100	1.11

- $T(n)/n^3$ is decreasing
- $f(n) = n^3$ looks like an upper bound
 - I.e., $T(n) \in O(n^3)$
 - If not tight, if the values tend towards 0, then $T(n) \in o(n^3)$
- We know that T grows no slower than a quadratic, and no faster than a cubic
- We *might* be able to improve one or both of those bounds

E.g. $-n^2 \log n$

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Consider $f_n = n^2 \log n$

n	$\frac{T(n)}{n^2 \log n}$
10	16.974
20	17.239
30	18.279
40	19.313
50	20.274
60	21.162
70	21.986
80	22.755
90	23.477
100	24.159

■ $\frac{T(n)}{n^2 \log n}$ also seems to be increasing

■ So, we have a new (better) lower bound

■ $T(n) \in \Omega(n^2 \log n)$

Let's try moving up a bit more:

E.g. – $n^{2.3}$

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Consider $f_n = n^{2.3}$

n	$T(n)/n^{2.3}$
10	19.589
20	21.024
30	22.411
40	23.558
50	24.528
60	25.369
70	26.112
80	26.781
90	27.388
100	27.947

- I'm comfortable saying $n^{2.3}$ is a lower bound
 - $T(n) \in \Omega(n^{2.3})$

Let's try moving up a bit more:

E.g. $-n^{2.5}$

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Consider $f_n = n^{2.5}$

n	$T(n)/n^{2.5}$
10	12.360
20	11.548
30	11.351
40	11.265
50	11.217
60	11.186
70	11.164
80	11.148
90	11.136
100	11.126

- This looks like an upper bound
 - $T_n \in O(n^2\sqrt{n})$
- Is it tight?
- Let's try something a little lower:

E.g. – $n^{2.4}$

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Consider $f_n = n^{2.4}$

n	$T(n)/n^{2.4}$
10	15.560
20	15.581
30	15.949
40	16.290
50	16.587
60	16.845
70	17.074
80	17.279
90	17.464
100	17.633

■ Increasing, so, lower bound

■ $T_n \in \Omega(n^{2.4})$

E.g. – Conclusion

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- We have T_n bound below by $n^{2.4}$ and bound above by $n^{2.5}$
- We maybe didn't find it exactly, but we have a very good idea how this algorithm grows
- Only push in each direction while you're comfortable w/the data
- None of this is proof
 - We need to choose input size sufficiently large to get past lower-order terms
 - No way to *know*
 - But, a program, on a given computer, has a practical upper limit on input size

Comparing Functions, Small n

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- You are generally not going to be able to find tight bounds, using the method above
- E.g., you can show that $3n \ln n \in o(n^{1.1})$
 - Hint: Use L'Hopital's Rule
- However $3n \ln n > n^{1.1}, \forall 3 < n < 1.2 \times 10^{16}$
- $\frac{3n \ln n}{n^{1.1}}$ will be increasing until around 66,000
- Remember, there's no such thing as a big number
- So, don't slice powers of n too finely

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh

Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Notes on Scale

Leading Zeros

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

Remember your Physics or Chemistry lab lessons

- As you divide by larger magnitudes, your data pick up leading zeros
 - These are *not* significant digits
 - Do not make a statement such as “There is only a difference of 0.00001”
 - This is meaningless, neither large nor small
- Scale the column:

n	$T(n)/n^2$	$T(n)/n^2(*10000)$
10000	0.000238	2.380
20000	0.0002576	2.576
30000	0.0002691	2.691
40000	0.0002764	2.764
50000	0.0002829	2.829
60000	0.0002856	2.856

Scaling, Scientific Notation

Run Time &
Performance

Kurt Schmidt

Intro

Profiling Code

Run-time
Analysis

Big-Oh
Ranking of Functions

Timing
Programs

Program
Growth

Notes on
Scale

- The scale is irrelevant
 - As long as it's consistent for a given column
- You can use scientific notation
 - Keep the exponent the same

n	$T(n)/f(n)$
10000	2.384 e-8
20000	1.288 e-8
30000	0.8970e-8
40000	0.6917e-8
50000	0.5659e-8
60000	0.4762e-8